

OSIREON - Core Development Guidance

Overview

This document provides the foundational architecture and development guidelines for the Osireon Core System. The goal is to build a modular, scalable, and transparent backend that connects seamlessly with the frontend and is ready to support plug-and-play applications, modules, and simulations. The system will initially be developed with **Italy** as the first target country, with plans to expand internationally.

1. Objectives

- **Modularity:** The system must be built around independent modules with clear APIs.
 - **Transparency:** Every process must be traceable and auditable.
 - **Scalability:** Must support future integrations and increased computational load.
 - **Security:** Built-in protection for data integrity, privacy, and access control.
 - **Interoperability:** Can connect with various LLMs, frontends, external datasets, and applications.
 - **Multinational Readiness:** Although initially focused on Italy, the system must support configuration and extension for other countries.
-

2. Tech Stack (Recommended)

- **Backend:** Python (FastAPI preferred for modular endpoints)
- **Database:** PostgreSQL (Relational), Redis (for caching and queueing)
- **Queueing:** Celery or FastAPI BackgroundTasks

- **Containerization:** Docker + Docker Compose
 - **Orchestration:** Kubernetes (future scaling)
 - **Frontend Bridge:** REST API (JSON), ready for GraphQL/WebSocket layer
 - **Authentication:** JWT-based access (to be upgraded with OAuth2/Keycloak in v2)
 - **LLM Interface:** OpenAI/LLM agnostic connector via API wrapper
-

3. Folder Structure (Initial Proposal)

```
osireon-core/
└── app/
    ├── main.py          # Entry point
    └── core/
        ├── engine.py     # Core logic engine
        │   ├── engine.py   # Main reasoning/processing loop
        │   └── decision_tree.py # Core decision/simulation structure
        ├── modules/       # Plug-in logic modules (LLM, economy, politics)
        ├── agents/         # Independent reasoning agents (Analyst, Critic, Historian, etc.)
        ├── adapters/       # Data source adapters (Eurostat, OWID, etc.)
        ├── api/            # REST API endpoints
        ├── db/              # DB models and interactions
        ├── services/        # Services (auth, logging, mail, utils)
        ├── schemas/         # Pydantic schemas
        ├── config.py
        └── dependencies.py
    ├── diagrams/        # Architecture and agent diagrams
    └── tests/
        └── Dockerfile
        └── docker-compose.yml
    └── README.md
```

4. Core Features (MVP Phase)

4.1. System Bootstrap

- Launches the API, database, and logger

- Loads configuration (ENV-based)

4.2. LLM Connector Module

- Standard interface to plug any LLM (e.g., OpenAI, Claude, open-source)
- Logs all prompts and responses
- Configurable models per task type

4.3. Policy Simulation Engine (Minimal)

- Receives inputs (JSON)
- Processes logic via a configurable reasoning tree
- Returns simulated impact or response

4.4. Modular Input-Output Layer

- Accepts input from:
 - Web frontend
 - Admin backend
 - External apps via API key
- Responds with JSON packages (logs, results, metadata)

4.5. Agent-Based Execution Model

- Each task is assigned to one or more **Agents** with specific skills:
 - AnalystAgent: runs simulations and evaluations
 - CriticAgent: provides counterarguments and detects inconsistencies
 - HistorianAgent: checks precedents and timeline impact

- Agents communicate via internal task queue or event bus
- Can be extended with future Agents (e.g., LegalAgent, EthicsAgent)

4.6. Data Adapter Layer

- Each data source is accessed via a standard adapter:

```
class BaseAdapter:
    def fetch(self, query): ...
    def normalize(self, raw): ...
```

- Adapters available:
 - `italy_istat_adapter.py`
 - `eurostat_adapter.py`
 - `owid_adapter.py`
 - `mock_adapter.py`

5. API Specification (Draft)

Base URL: `/api/v1`

POST /simulate

- Description: Runs a simulation based on provided policy input
- Payload:

```
{
  "country": "Italy",
  "domain": "Tax Reform",
  "proposals": ["Flat tax 20%", "Universal Basic Income"],
```

```
        "constraints": ["No increase in national debt"]  
    }  
  
    ● Response:
```

```
{  
    "results": {...},  
    "reasoning_tree": {...},  
    "ilm_output": "...",  
    "meta": {  
        "execution_time": 2.4,  
        "modules_used": ["economy", "social"]  
    }  
}
```

6. Modules & Extensibility

Each module (e.g., economy, climate, education) is a Python class/function set with defined:

- `register()` method
- `execute(input)` method
- `metadata` and `version`

Modules are auto-loaded from `/modules` folder using a loader script. Each module also contains:

- Status: `stable`, `experimental`, or `deprecated`
 - Compatible countries
 - Required data inputs
-

7. Frontend Interface Hook

- REST API returns ready-to-render JSON
 - Standard data format for charts, timelines, recommendations
 - Add WebSocket hooks for live simulation and stream response (future)
-

8. Dev Notes

- Codebase must be well-documented and fully typed
 - Use environment variables for all secrets and model keys
 - Set up logging and error tracking (e.g., Sentry integration optional)
 - All inputs/outputs must be logged in the DB for transparency
 - Include initial unit test framework with coverage tracking
-

9. Future Goals (not MVP but planned)

- User management system
- Admin panel for simulation logs
- Full pipeline of scenario comparison and rating
- External API for contributors (public endpoint + rate limiting)
- Governance module integration
- Agent-to-agent communication layer
- Plugin interface for civic tech tools (Decidim, vTaiwan, etc.)
- Versioning dashboard for modules

10. Licensing & Ethics

- Core must follow AGPL v3 license
 - Commercial/Pro modules can be added under dual-licensing system
 - Code of ethics check in every module (auto-log if violated)
-

11. Diagrams & Visual Architecture (To Be Added)

- High-level system architecture (backend ↔ LLM ↔ frontend ↔ DB)
- Internal flow diagram (input → agents → modules → output)
- Modular architecture map
- Agent hierarchy and collaboration model

All diagrams will be saved in `/diagrams` in `.drawio` and `.png` formats.

Summary for Dev

Start by building:

- The API base (FastAPI)
- The logic engine scaffold with module loading
- A working simulation endpoint
- LLM wrapper for OpenAI
- Logs into PostgreSQL
- Mock modules: `economy_mock`, `ethics_mock`, and `italy_stats`

12. Extended Architecture & Advanced Features

The following elements expand the Osireon Core into a fully verifiable, ethics-driven, and community-integrated system.

12.1. Ethics Engine Codified on Blockchain

- A module within the Core Engine that verifies every decision against an immutable **Code of Ethics**.
- Each rule is versioned, hashed, and optionally stored on a blockchain (e.g., Ethereum/Arbitrum).
- The engine signs each simulation and stores the signature hash for later auditing.
- Verifier nodes (optional) can independently confirm that simulations respected all ethical rules.

Files and Functions:

/ethics/

```
|—— ruleset_v1.json  
|—— validator.py      # Applies rules  
|—— blockchain_client.py # Interface with smart contracts or decentralized log
```

12.2. License & Module Validation Layer

- A LicenseManager validates modules using:
 - License keys (for commercial modules)
 - Signature hash verification for integrity
- Open source modules are tagged as AGPL by default.
- Modules are loaded only if license and checksum pass.

Files:

/licenses/

```
|—— manager.py  
|—— agpl_checker.py  
|—— commercial_checker.py
```

12.3. Contributor & DAO Interface

- Community members can submit new policies, modules, or feedback.
- Contributions go through a ProposalReviewFlow, optionally governed by:
 - Moderators
 - DAO voting (future)
- Approved contributions are logged and versioned.

Files:

/contributors/

```
|—— portal.py  
|—— proposal_flow.py  
|—— dao_interface.py (future)
```

12.4. Public Audit Logging

- All simulations and decisions are saved to:
 - Local DB (PostgreSQL)
 - Decentralized log (e.g., Arweave/IPFS) for transparency

Logged data:

- Simulation ID
- Input, timestamp
- Output, reasoning
- Ethics hash & validator status
- Modules used & agents involved

12.5. Real-Time Simulation Layer (WebSocket API)

- For live dashboards and interactions.
- Streaming data (chunked responses) from agent evaluations and LLM reasoning.

Endpoint example:

/ws/simulation/live?id=...

12.6. Country-Specific Configurations

- Each country node includes:
 - Language pack
 - Local data adapters

- Localized modules (legal, economic)
- Country configs can override default modules

Structure:

/countries/

```
|   └── italy/
|       ├── modules/
|       ├── adapters/
|       └── translations.json
|
└── france/
    └── usa/
```

Next Steps

- Implement core components from section 12 in parallel with MVP.
 - Ethics and License layers can begin as mock checkers with logs.
 - Community portal can be launched with proposal intake only, DAO later.
-

This extended architecture completes the vision of Osireon as a transparent, ethical, decentralized decision engine for the public.

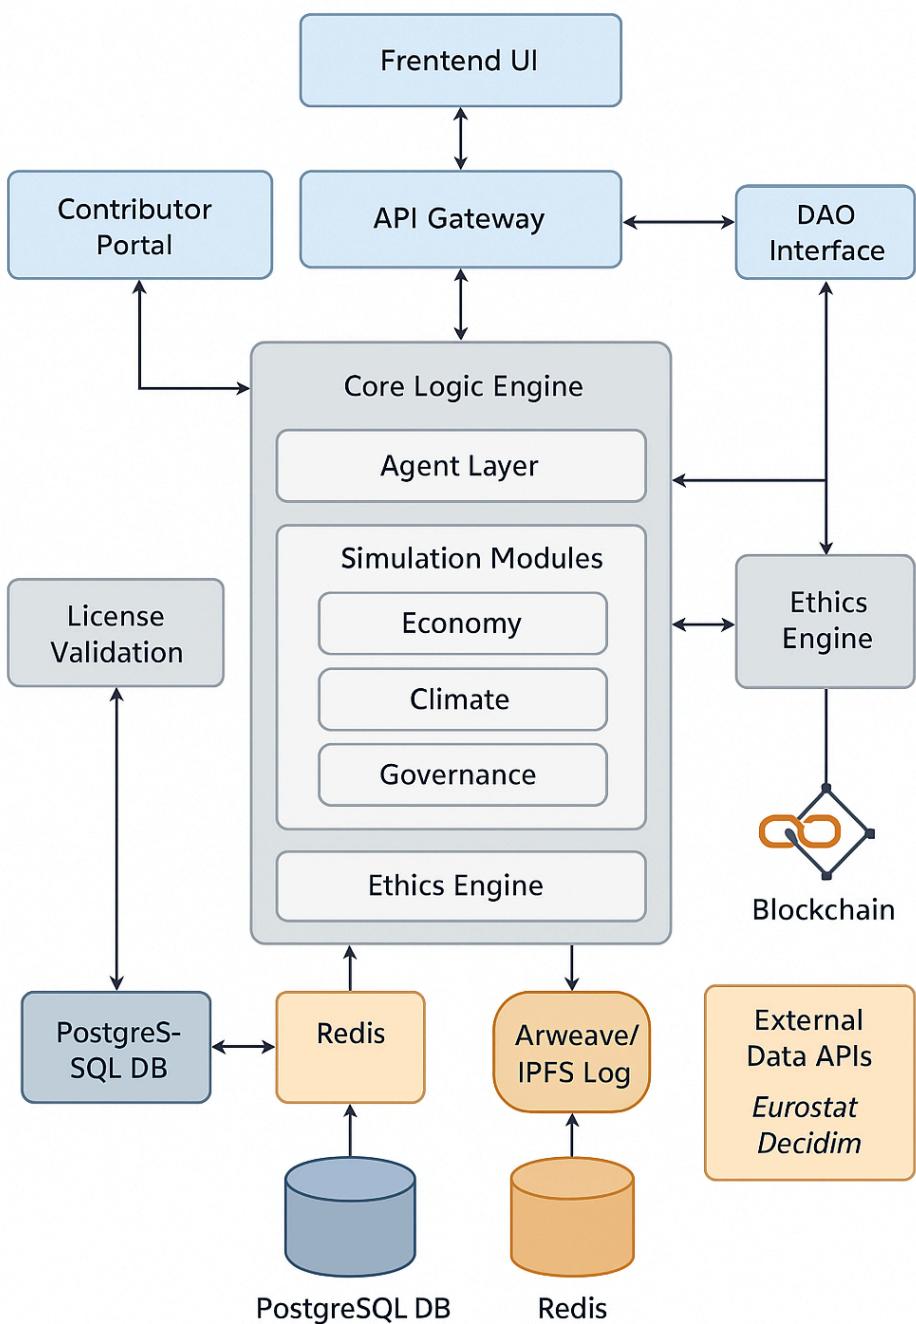


OSIREON System Diagrams – Overview & Prompts

Each of the following diagrams represents a critical layer of the Osireon Core Architecture. They help developers, contributors, and users visually understand how data flows, agents interact, and modules connect.

1. High-Level System Architecture

🎯 **Purpose:** Show the main components of Osireon and how they communicate: API, Core Engine, Modules, Agents, Frontend, LLM, and DB.



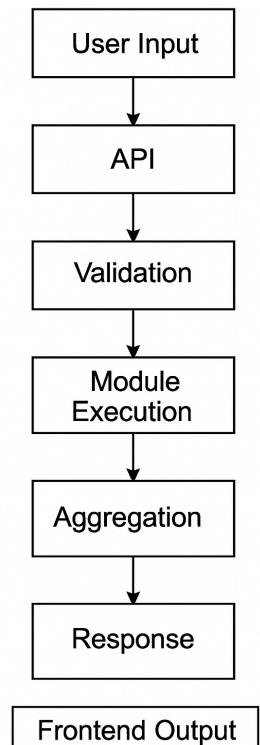
Osireon Civic tech AI eciscon1.5+

💬 Diagram Comment:

This shows the entire Osireon ecosystem at a glance. Each box is an independent service or layer that can evolve separately while staying interoperable with the whole system.

2. Simulation Data Flow

🎯 **Purpose:** Visualize the internal path of a simulation request, from input to final response.



💬 **Diagram Comment:**

This diagram reveals how Osireon transforms a single user input into a rich, multi-layered simulation using structured reasoning and AI assistance.

3. Modular Agent & Module Architecture

🎯 **Purpose:** Show how agents and modules are dynamically loaded, registered, and invoked.

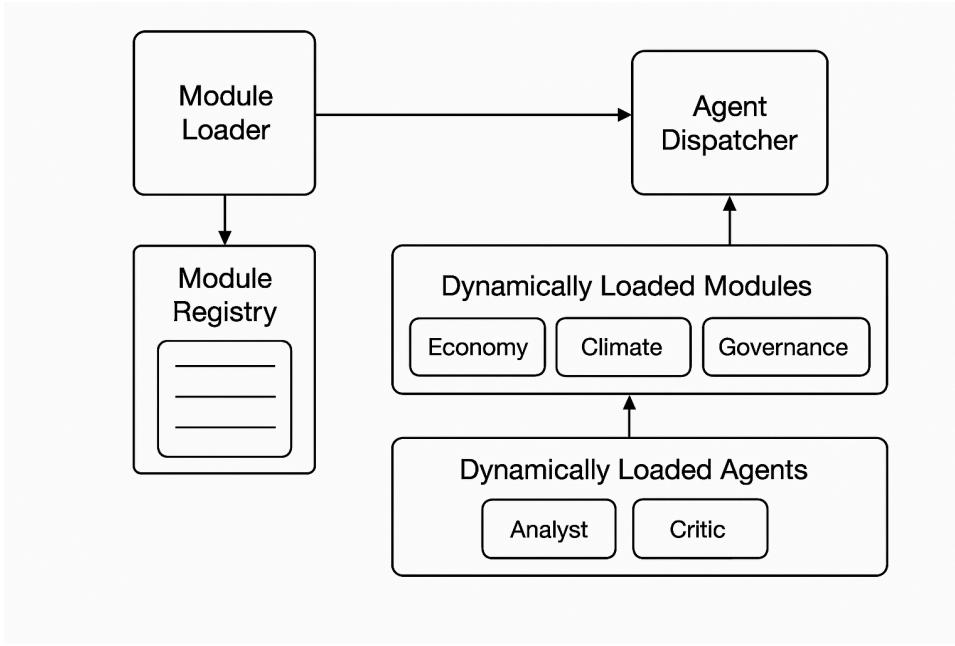
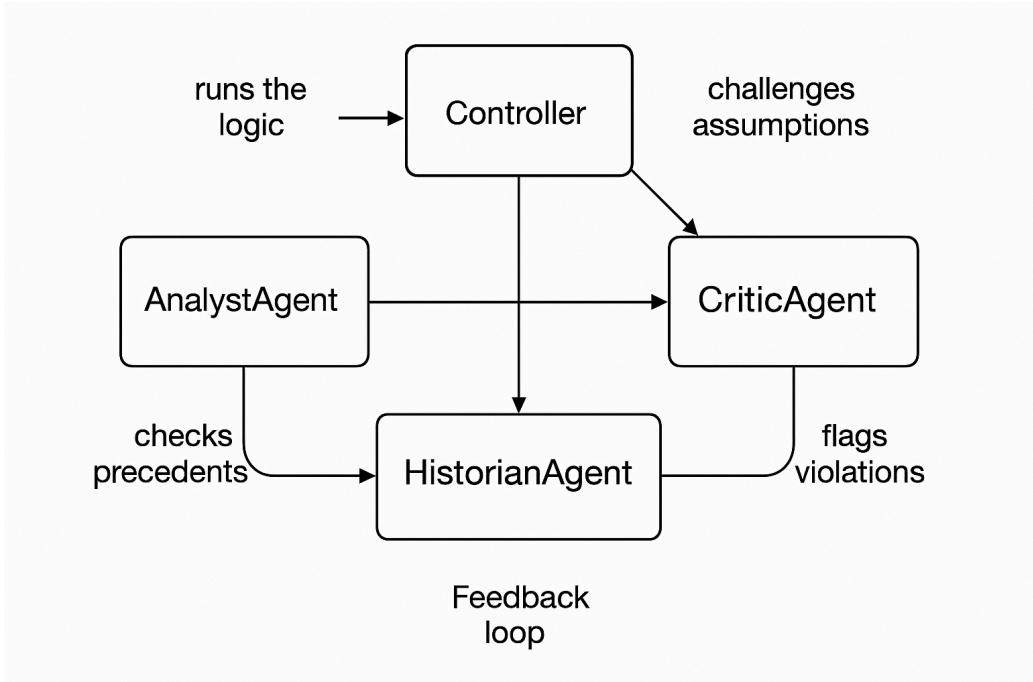


Diagram Comment:

This is the brain of Osireon. Everything is modular: agents act as specialized minds, modules as policy knowledge domains. The loader connects them in real time based on context.

4. Agent Collaboration & Interaction Map

Purpose: Display how agents collaborate, correct each other, and converge on a final output.



 **Diagram Comment:**

Osireon simulates reasoned debate. Each agent has a defined purpose and checks the others, ensuring that proposals are thoroughly tested before becoming recommendations.

5. External Data Adapter Layer

 **Purpose:** Show how Osireon connects to external databases and civic tech APIs via adapters.

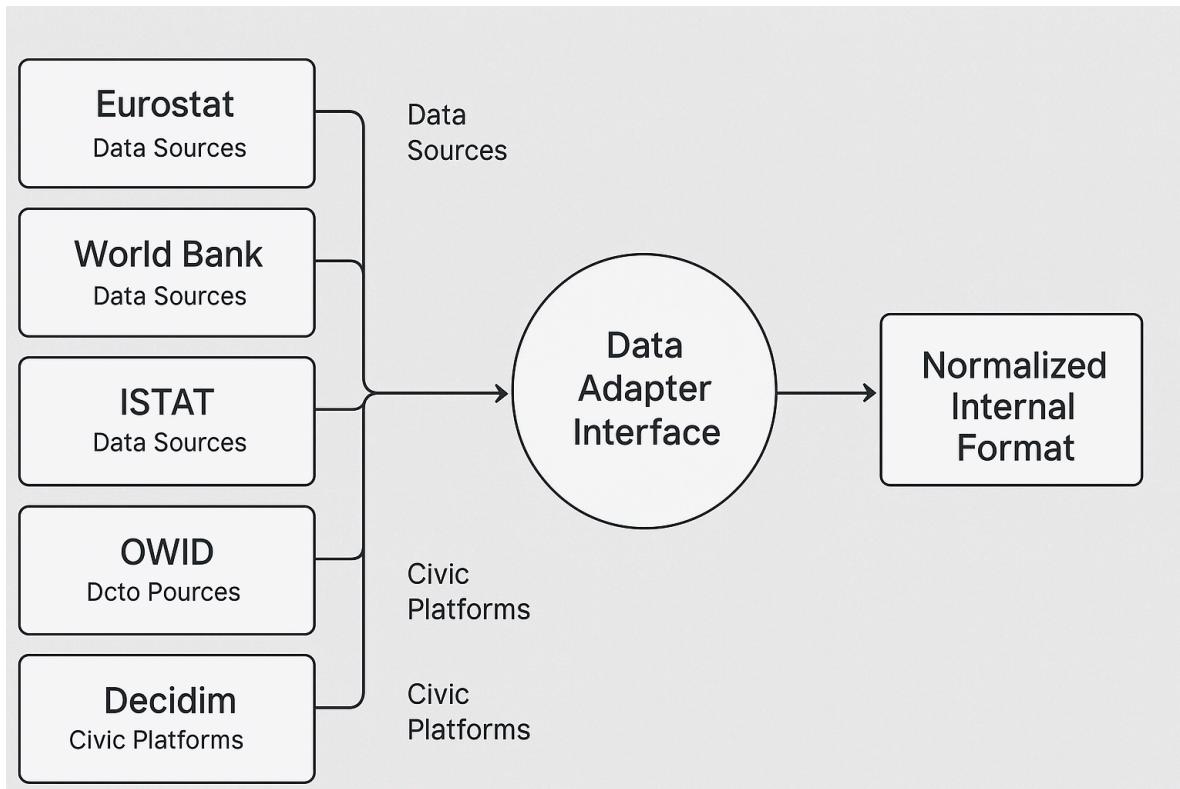
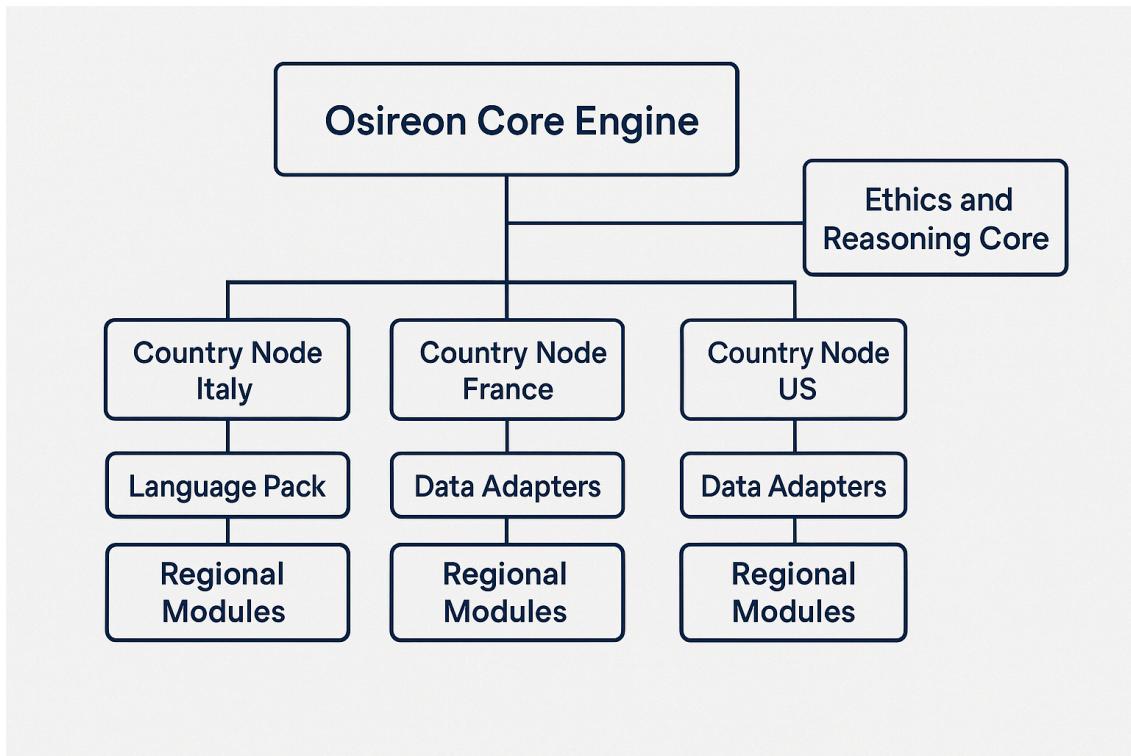


Diagram Comment:

Osireon doesn't guess – it grounds its logic in real-world facts. This adapter layer acts like a translator between diverse external datasets and the internal logic engine.

6. Scalable Multi-Country Configuration

Purpose: Display how Osireon can support multiple countries via independent configurations.



💬 **Diagram Comment:**

Osireon is designed for global impact. Every country can have a custom layer on top of a shared reasoning core—allowing local adaptation without compromising global coherence.

7. Ethics Engine & Blockchain Validation Flow

🎯 **Objective:**

To visualize how policy simulations are validated against a codified code of ethics and recorded on the blockchain for immutability and transparency.

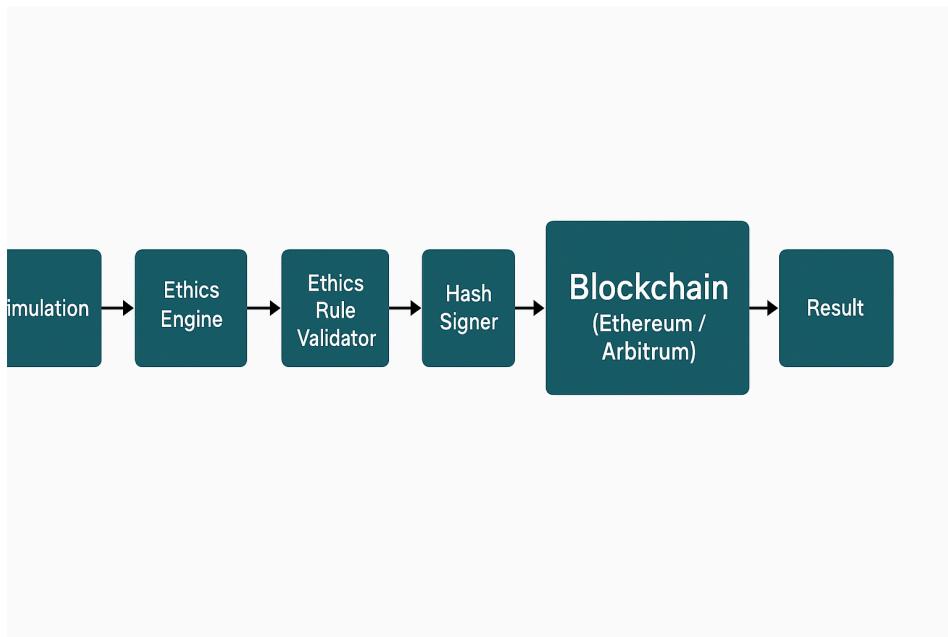


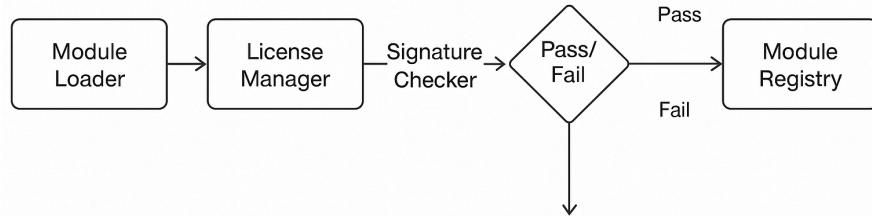
Diagram Comment:

Every simulation in Osireon is checked against an ethics ruleset. Once validated, the result is cryptographically signed and optionally stored on-chain to ensure public accountability and traceability.

8. License and Module Integrity Verification Flow

Objective:

To illustrate the process that verifies the license and integrity of each module before it is loaded into the system.



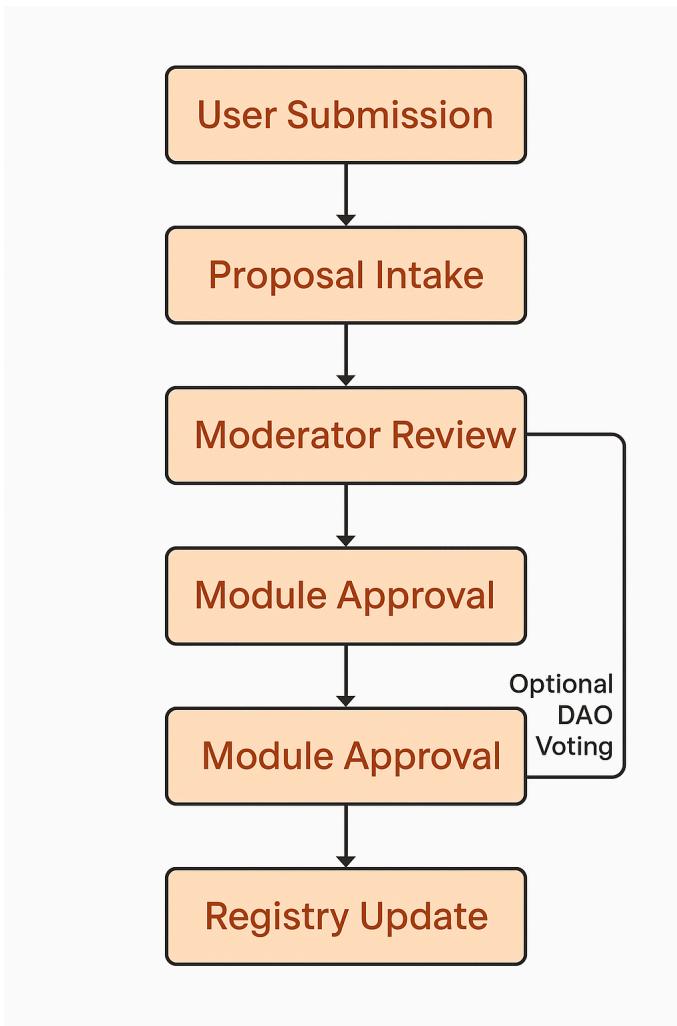
💬 **Diagram Comment:**

This security layer protects Osireon's ecosystem from tampered or unauthorized modules. Only signed and licensed modules are allowed to interact with the core engine.

9. Contributor Proposal Flow & DAO Governance

🎯 **Objective:**

To show how new modules, simulations, or ideas submitted by users are reviewed, optionally voted by a DAO, and then added to the system.



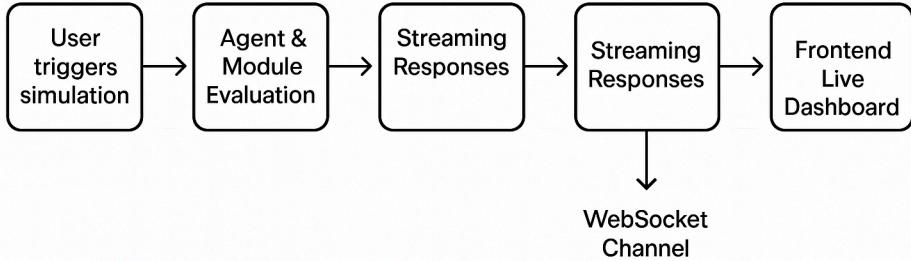
💬 **Diagram Comment:**

This community flow ensures that innovation is decentralized, reviewed fairly, and managed transparently through democratic or delegated processes.

10. Real-Time Simulation & WebSocket Layer

🎯 **Objective:**

To demonstrate how a user-triggered simulation streams results live to the frontend using WebSocket for real-time feedback.



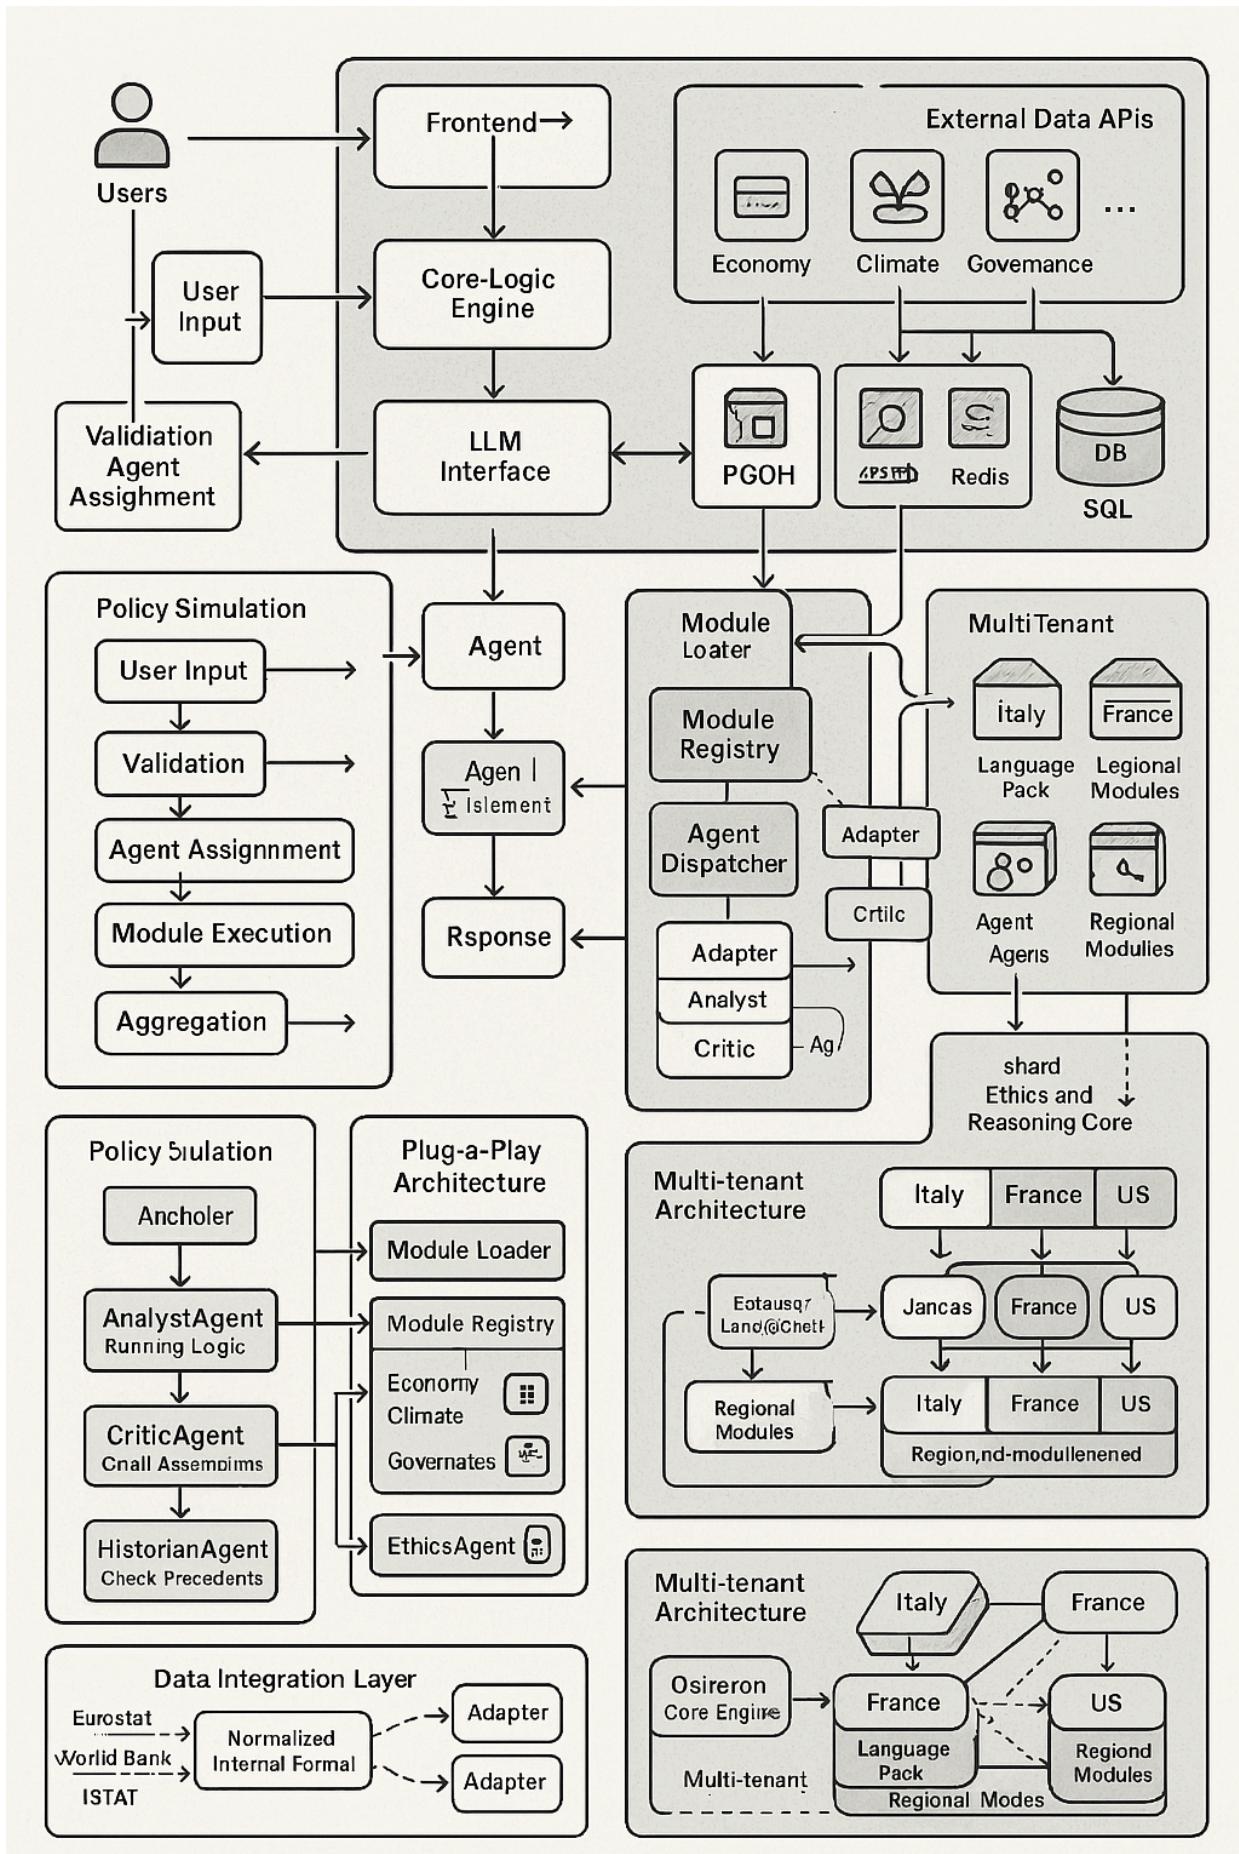
💬 Diagram Comment:

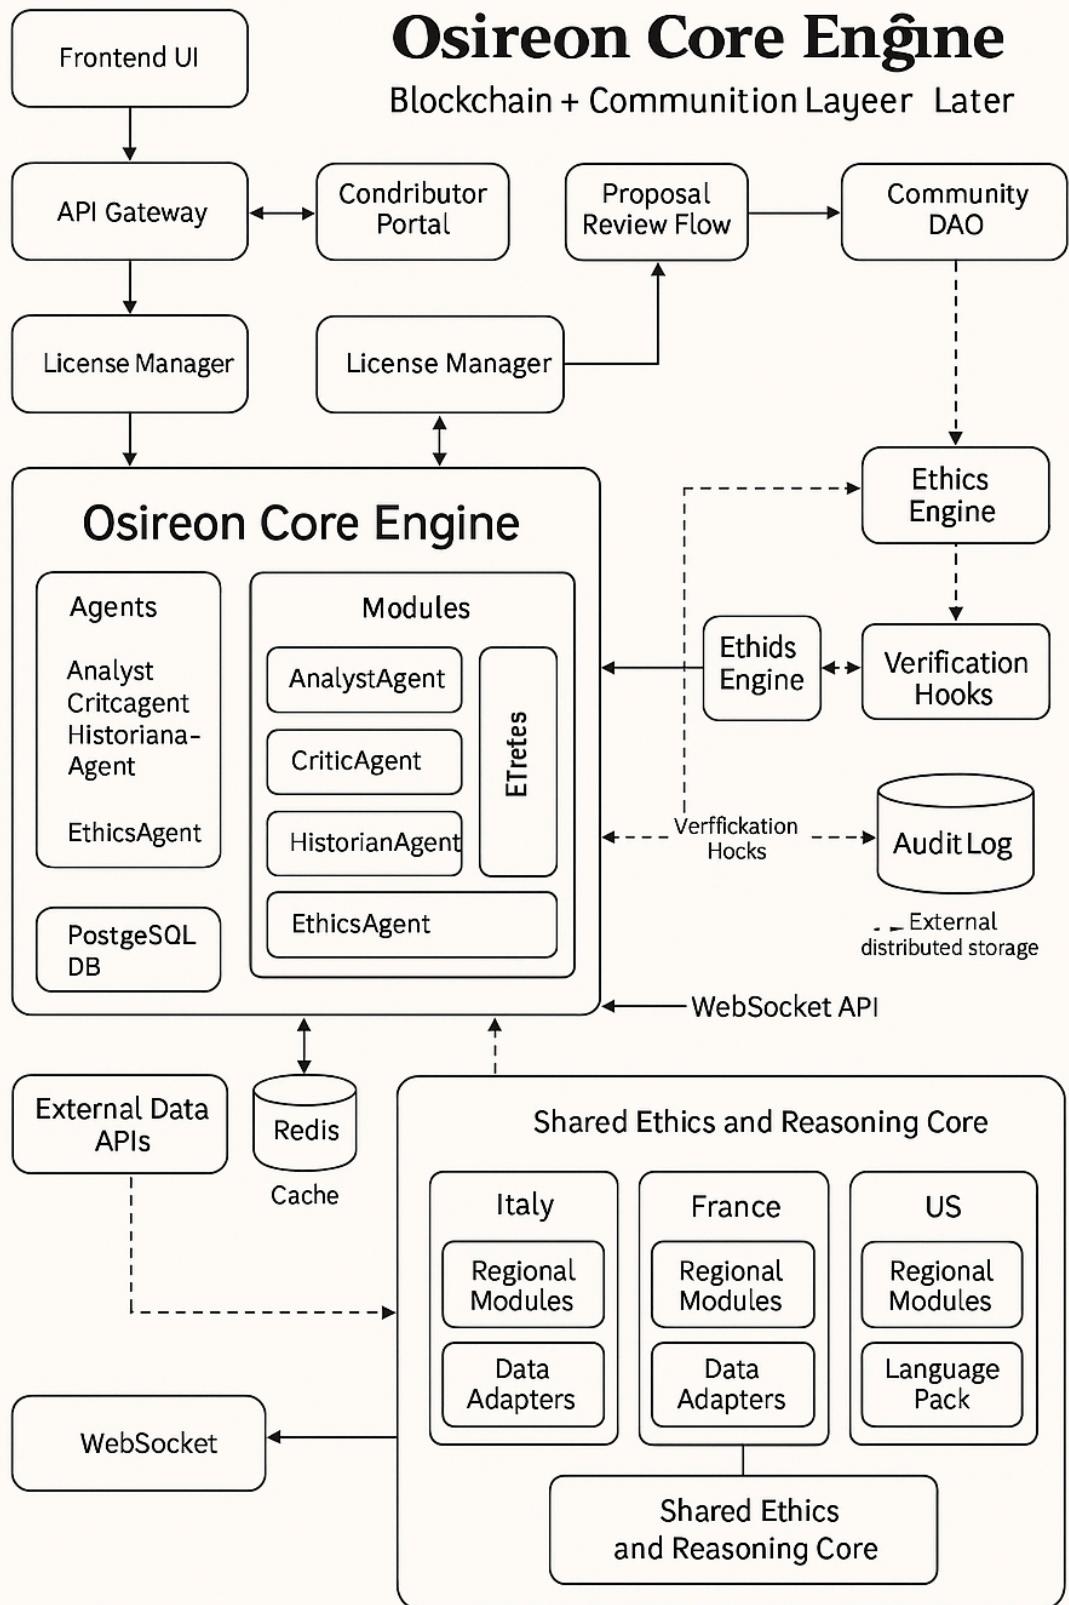
With WebSocket integration, Osireon becomes a live civic lab. Users can visualize how decisions evolve in real time as agents and modules process them collaboratively.

Osireon Full System Architecture

🎯 Objective:

To provide a comprehensive overview of the Osireon v1.5+ architecture, including new advanced components like the Ethics Engine, Licensing System, DAO, Public Audit Log, and WebSocket support.





 **Diagram Comment:**

This is the complete system view of Osireon. It integrates community participation, ethical enforcement, decentralized governance, and real-time capabilities — all connected to a modular and auditable simulation engine.

Questions?

Ping Giorgio or submit ideas/issues directly into the GitHub repo.
